

Understanding and Detecting Concurrency Attacks

Paper #82

The University of Hong Kong
@cs.hku.hk

Abstract

Just like sequential bugs lead to attacks, concurrency bugs also lead to concurrency attacks.

1. Introduction

Multi-threaded program is hard to be correct. Concurrency bugs are common in modern multi-threaded programs including atomic violation, ..., and especially data race[13, 14, 22, 23]. Extant work well explores interleaving that causes concurrency bugs, and efficiently detects explicit concurrency bugs that direct to severe consequences such as execution order violation, wrong output and program crash [13, 15, 18, 19, 22, 23].

Recent studies[17, 21] show rise of concerns about *concurrency attacks*. By triggering concurrency bugs, hackers may leverage the corrupted memory to conduct attacks including privilege escalations[7, 9], hijacking code execution[8], bypassing security checks[3–5], and breaking database integrity[17]. These vulnerabilities are often hidden in concurrency bugs and implicit regarding program behaviors. For example, a privilege escalation attack may cause no outrageous effect, but possess a permanent security hole hiding in the system. Also, despite the inputs that induce concurrency bugs, concurrency attacks may need other crafted inputs for exploitation of the vulnerabilities.

Unfortunately, although great progress has been made, there does not exist a general model and practical tools for understanding and detecting concurrency attacks. Our study over known concurrency bugs[1, 2] shows that even concurrency bugs have been successfully detected and reported, professionals may still lack knowledge about how severe consequences these bugs may cause. For instance, *apache-25520*[1] has been reported over years and well studied by researchers[14]. We are the first to exploit a new heap over-

flow attack leveraging on this bug and break the HTML integrity. Knowledge and detection of concurrency attacks is of crucial importance.

We studied 26 attacks and find two major challenges for detecting concurrency attacks. First, concurrency attacks may be implicit and hidden in common concurrency bugs. Existing concurrent bug detectors may not analyze whether a concurrency bug is vulnerable and exploitable. ConMem[23] first propose to consider concurrency bugs that may cause severe consequences (e.g. program crash), but our observation shows that explicit error (e.g. program crash) is not necessary for concurrency attacks. ConSeq[22] care about severe consequences and do intra-procedure analysis to help diagnose concurrency bugs. However, propagation path Worse still, crash bugs may even lead to more severe vulnerabilities. In *CVE-2017-7533* conducted by our team, although the data race primarily causes kernel crash, we crafted the input and successfully conduct a privilege escalation attack without crashing the kernel. Current concurrency bug detecting tools are not designed to analyze this kind of latent vulnerabilities, and hence may direct wrong level of warnings towards the bugs.

Second, extant work ignores indicating extra inputs to conduct concurrency attacks. A concurrency bug may become much more vulnerable when attackers 1.craft inputs that trigger the bug; 2.employee another input running on victim threads to construct their attacks. In *CVE-2017-7533*, we do not only leverage two crafted inputs running on two threads to trigger a data race and construct kernel heap overflow, but also require another inputs running on *victim thread* to lay the target structure on the same heap. By corrupting the target structure, we finally achieve arbitrary code execution and get a root shell. Automatically indicating the inputs that construct concurrency attacks would be of vital helpful for developers to better understand the latent vulnerabilities.

To address the two challenges, we present a general model (§3.1) for understanding concurrency attacks. The model breaks down concurrency attacks into three stages: bug happening, bug-to-attack propagation, and attack happening. In this model, a concurrency bugs is triggered by two

Leveraging this model, we designed a two-phase framework (§3.2), XXX, for detecting concurrency attacks.

The first phase is *concurrency analyzer* to analyze bug-to-attack propagations. Our study found that most vulnerable races are already included in the race detectors reports, and concurrency attacks sites are often explicit in program code. Therefore, we can perform static analysis on only the data and control flow propagations between the bug reports and the potential attack sites, then we can collect relevant call stacks and branch statements as the potentially vulnerable input hints.

The second phase is *concurrency fuzzer* to ...

We implemented XXX on Linux, supporting both user space and kernel space attack detection.

We evaluated OWL on 6 diverse, widely used programs, including Apache, Chrome, Libsafe, Linux kernel, MySQL, and SSDB. OWLs benign schedule hints and runtime verifiers reduced 94.3% of the race reports, and it did not miss the evaluated concurrency attacks. With the greatly reduced reports, OWLs vulnerable input hints helped us identify subtle vulnerable inputs, leading to the detection of 7 known concurrency attacks as well as 3 previous unknown, severe ones in SSDB and Apache. The analysis performance of OWL was reasonable for in-house testing.

This paper makes two major contributions:

1. **A general model for understanding concurrency attacks.** This model explains most concurrency attacks in wild and providing two major direction for detecting concurrency attacks.
2. **A practical concurrency attack detection tool and its implementation, XXX.** XXX can easily employ existing concurrent bug detectors and vulnerability analyzer to improve the accuracy and ??? of detection.

The rest of this paper is structured as follows. §2 introduces the background of concurrency attacks. §3 gives an overview on the concurrency attack model and architecture of XXX. §??...

2. Background

2.1 Concurrency Bug

2.2 Concurrency Attack

Extant studies [17, 21] concurrency attacks into .. concurrency bugs and made three major findings on concurrency attacks. First, concurrency attacks make much more severe threats: 35 of the bugs can corrupt critical memory and cause four types of severe security consequences, including privilege escalations [7, 10], malicious code injections [8], and bypassing security authentications [4, 3, 5]. Second, concurrency bugs and attacks can often be easily triggered via subtle program inputs. For instance, attackers can use inputs to control the physical timings of disk IO and program loops and trigger concurrency bugs with a small number of re-executions. Third, compared to traditional TOCTOU attacks, which stem from corrupted file accesses, handling concurrency attacks is much more difficult because they stem from corrupted, miscellaneous memory accesses.




Figure 1: Concurrency Attack Model

2.3 Bug-to-Attack Propagation

3. Overview

We introduce a model that explains most concurrency attacks we studied. Leveraging the model, we design a framework XXX to detect concurrency attacks. This section gives an overview of the model and architecture of XXX.

3.1 Concurrency Attack Model

3.2 XXX's Architecture

4. XXX's Framework

5. Implementation

6. Discussions

7. Evaluation

We evaluated XXX on 6 widely used C/C++ programs, including three server applications (Apache [11] web server, MySQL [10] database server, and SSDB [16] key-value store server), one library (Libsafe [12]), the 4.11.9 Linux kernel, and one web browser (Chrome). We used the programs' common performance benchmarks as workloads. Our evaluation was done on XXX.

We focused our evaluation on four key questions:

1. How many false reports from concurrency error detection tools can XXX reduce (§??)?
2. How many potential victim threads can XXX indicate (§3)?
3. Can XXX detect known concurrency attacks in the real-world (§??)?
4. Can XXX detect previously unknown concurrency attacks in the real-world (§??)?

8. Related Work

9. Conclusion

References

- [1] Apache bug 25520. https://bz.apache.org/bugzilla/show_bug.cgi?id=25520.

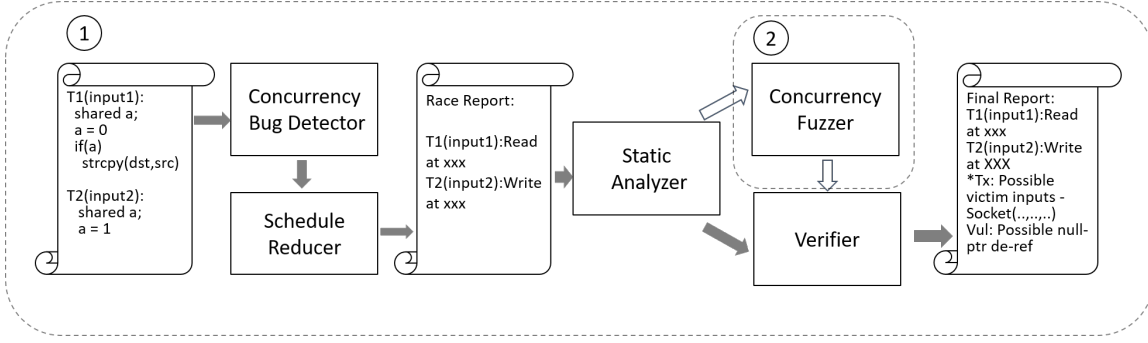


Figure 2: Architecture

Name	LoC	# r.r.	# XXX's r.	# atks	# atks found
Apache	290K	715	10	3	3
Chrome	3.4M	1715	115	1	1
Libsafe	3.4K	3	3	1	1
Linux	2.8M	24641	34	2	2
MySQL	1.5M	1123	16	2	2
SSDB	67K	12	2	1	1
Total	-	-	-	-	-

Table 1: XXX race report reduction and concurrency attack detection results. Description

Name	Type	# Fed Inputs	# Victim Inputs	# atks
Apache	apr_palloc	-	-	-
Linux	kmalloc	180	6	1
Total	-	-	-	-

Table 2: XXX's Concurrency Fuzzer. Description

- [2] Apache bug 46215. https://bz.apache.org/bugzilla/show_bug.cgi?id=46215.
- [3] CVE-2008-0034. <http://www.cvedetails.com/cve/CVE-2008-0034/>.
- [4] CVE-2010-0923. <http://www.cvedetails.com/cve/CVE-2010-0923>.
- [5] CVE-2010-1754. <http://www.cvedetails.com/cve/CVE-2010-1754/>.
- [6] CVE-2017-7533. <https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2017-7533>.
- [7] Linux kernel bug on uselib(). <http://osvdb.org/show/osvdb/12791>.
- [8] MSIE javaprxy.dll COM object exploit. <http://www.exploit-db.com/exploits/1079/>.
- [9] Mysql bug 14747. <https://bugs.mysql.com/bug.php?id=14747>.
- [10] MySQL Database. <http://www.mysql.com/>, 2014.
- [11] Apache web server. <http://www.apache.org>, 2012.
- [12] Libsafe. <http://directory.fsf.org/wiki/Libsafe>.
- [13] S. Lu, S. Park, C. Hu, X. Ma, W. Jiang, Z. Li, R. A. Popa, and Y. Zhou. Muvi: automatically inferring multi-variable access correlations and detecting related semantic and concurrency bugs. In *Proceedings of the 21st ACM Symposium on Operating Systems Principles (SOSP '07)*, pages 103–116, 2007.

- [14] S. Lu, S. Park, E. Seo, and Y. Zhou. Learning from mistakes: a comprehensive study on real world concurrency bug characteristics. In *Thirteenth International Conference on Architecture Support for Programming Languages and Operating Systems (ASPLOS '08)*, pages 329–339, Mar. 2008.
- [15] N. Nethercote and J. Seward. Valgrind: a framework for heavyweight dynamic binary instrumentation. In *Proceedings of the ACM SIGPLAN 2007 Conference on Programming Language Design and Implementation (PLDI '07)*, pages 89–100, June 2007.
- [16] ssdb.io/.
- [17] P. B. Todd Warszawski. Acidrain: Concurrency-related attacks on database-backed web applications. In *Proceedings of the 2017 ACM SIGMOD International Conference on Management of Data*, pages 5–20. ACM, 2017.
- [18] Threadsanitizer. <https://code.google.com/p/data-race-test/wiki/ThreadSanitizer>, 2015.
- [19] Z. Wu, K. Lu, X. Wang, and X. Zhou. Collaborative technique for concurrency bug detection. *International Journal of Parallel Programming*, 43(2):260–285, 2015.
- [20] W. Xiong, S. Park, J. Zhang, Y. Zhou, and Z. Ma. Ad hoc synchronization considered harmful. In *Proceedings of the Ninth Symposium on Operating Systems Design and Implementation (OSDI '10)*, Oct. 2010.
- [21] J. Yang, A. Cui, S. Stolfo, and S. Sethumadhavan. Concurrency attacks. In *the Fourth USENIX Workshop on Hot Topics in Parallelism (HOTPAR '12)*, June 2012.
- [22] W. Zhang, J. Lim, R. Olichandran, J. Scherpelz, G. Jin, S. Lu, and T. Reps. ConSeq: detecting concurrency bugs through sequential errors. In *Sixteenth International Conference on Architecture Support for Programming Languages and Operating Systems (ASPLOS '11)*, pages 251–264, Mar. 2011.
- [23] W. Zhang, C. Sun, and S. Lu. ConMem: detecting severe concurrency bugs through an effect-oriented approach. In *Fifteenth International Conference on Architecture Support for Programming Languages and Operating Systems (ASPLOS '10)*, pages 179–192, Mar. 2010.