

Understanding and Detecting Concurrency Attacks

Paper #82

The University of Hong Kong
@cs.hku.hk

Abstract

1. Introduction

Multi-threaded program is hard to be correct. Concurrency bugs are common in modern multi-threaded programs[3, 4, 8, 9] including atomic violation, ..., and especially data race. Extant work well explore interleaving that cause concurrency bugs, and efficiently detect explicit concurrency bugs that direct to severe consequences such as execution order violation, wrong output and program crash.

Recent studies[5, 7] show rise of concern about *concurrency attacks*. Just like sequential bugs leads to attacks, concurrency bugs also lead to concurrency attacks. By triggering concurrency bugs and employing subtle inputs, hackers may leverage the corrupted memory to conduct attacks including privilege escalations[], hijacking code execution[], bypassing security checks[], and breaking database integrity[5]. These vulnerabilities often hide in large amount of concurrency bug reports. For example, bug information leveraged by a xxx attack is hidden in 1000 race reports produced by TSAN[6], a famous and widely used data race detector. Also, despite threads conducting concurrency bugs, behaviors of another thread may also be affected when data race bugs infecting shared memory (e.g.heap overflow)[1?].

Unfortunately, although great progress has been made to detect and replay severe bugs(e.g.ConMem[9]), extant work still lacks exploration of concurrency attacks from enormous concurrency bugs. Our study over several known concurrency bugs[1, 2] shows that even concurrency bugs have been successfully detected and reported, professionals may still lack knowledge about how severe consequences these bugs may cause. For instance, *apache-25520*[1] has been reported over years and well studied by researchers[4]. We are

the first to exploit a new heap overflow attack leveraging on this bug and break the HTML integrity.

We studied 26 attacks and find two major challenges for exploring concurrency attacks from concurrency bugs. First, concurrency attacks may be implicit and hidden in common concurrency bugs. Existing concurrent bug detectors focus on bug happening itself. ConMem[9] first propose to consider concurrency bugs that may cause severe consequences (e.g.program crash) and ConSeq[8] uses severe consequence report to help diagnose concurrency bugs. However, our observation shows that explicit error (e.g.program crash) is not necessary for concurrency attacks. Some concurrency bugs (e.g.xxx-xxx) may not cause program crash or interrupted, but attackers may still leverage them and conduct attacks with crafted input. Worse still, crash bugs may even lead to more severe vulnerabilities. In *CVE-2017-7533* conducted by our team, although the data race primarily causes kernel crash, we crafted the input and successfully conduct a privilege escalation attack without crashing the kernel. Current concurrency bug detecting tools are not designed to analyze this kind of latent vulnerabilities, and hence may direct wrong level of warnings towards the bugs.

Second, extant work ignores indicating the *victim thread* of concurrency attacks. A concurrency bug may become much more vulnerable when attackers employee another thread to construct their attacks. In *CVE-2017-7533*, we do not only leverage two threads to trigger a data race and construct kernel heap overflow, but also require another *victim thread* to lay the target structure on the same heap. By crafting inputs and corrupt the target structure, we finally achieve arbitrary code execution and get a root shell. Automatically indicating the victim thread would be of vital helpful for developers to better understand the latent vulnerabilities. For example, in the *Apache-25520* case, after knowing about information of victim thread, we successfully increased the severity of the bug and conducted an integrity violation.

To address the two challenges, we introduced a new model that explains most concurrency attacks we studied. The model breaks down concurrency attacks into three stages: bug happening, bug-to-attack propagation, and attack happening. Bug happening is where the concurrency bugs happened.

Leveraging this model, we designed a two-phase framework, XXX, for detecting concurrency attacks. The first phase is *concurrency analyzer* to .

The second phase is *concurrency fuzzer*? to

We implemented XXX on Linux, supporting both user space and kernel space attack detection.

We evaluated OWL on 6 diverse, widely used programs, including Apache, Chrome, Libsafe, Linux kernel, MySQL, and SSDB. OWLs benign schedule hints and runtime verifiers reduced 94.3% of the race reports, and it did not miss the evaluated concurrency attacks. With the greatly reduced reports, OWLs vulnerable input hints helped us identify subtle vulnerable inputs, leading to the detection of 7 known concurrency attacks as well as 3 previous unknown, severe ones in SSDB and Apache. The analysis performance of OWL was reasonable for in-house testing.

This paper makes two major contributions:

1. **A general model explains happening and exploiting of concurrency attacks.** This model explains most concurrency attacks in wild and providing two major direction for detecting concurrency attacks.
2. **A general concurrency attack detection framework and its implementation, XXX.** XXX can easily employ existing concurrent bug detectors and vulnerability analyzer to improve the accuracy and ??? of detection.

2. Background

2.1 Concurrency Attack

Extant studies [5, 7] category concurrency attacks into concurrency bugs and made three major findings on concurrency attacks. First, concurrency attacks make much more severe threats: 35 of the bugs can corrupt critical memory and cause four types of severe security consequences, including privilege escalations [7, 10], malicious code injections [8], and bypassing security authentications [4, 3, 5]. Second, concurrency bugs and attacks can often be easily triggered via subtle program inputs. For instance, attackers can use inputs to control the physical timings of disk IO and program loops and trigger concurrency bugs with a small number of re-executions. Third, compared to traditional TOCTOU attacks, which stem from corrupted file accesses, handling concurrency attacks is much more difficult because they stem from corrupted, miscellaneous memory accesses.

8. Conclusion

References

- [1] Apache bug 25520. https://bz.apache.org/bugzilla/show_bug.cgi?id=25520.
- [2] Apache bug 46215. https://bz.apache.org/bugzilla/show_bug.cgi?id=46215.
- [3] S. Lu, S. Park, C. Hu, X. Ma, W. Jiang, Z. Li, R. A. Popa, and Y. Zhou. Muvi: automatically inferring multi-variable access correlations and detecting related semantic and concurrency bugs. In *Proceedings of the 21st ACM Symposium on Operating Systems Principles (SOSP '07)*, pages 103–116, 2007.
- [4] S. Lu, S. Park, E. Seo, and Y. Zhou. Learning from mistakes: a comprehensive study on real world concurrency bug characteristics. In *Thirteenth International Conference on Architecture Support for Programming Languages and Operating Systems (ASPLOS '08)*, pages 329–339, Mar. 2008.
- [5] P. B. Todd Warszawski. Acidrain: Concurrency-related attacks on database-backed web applications. In *Proceedings of the 2017 ACM SIGMOD International Conference on Management of Data*, pages 5–20. ACM, 2017.
- [6] Threadsanitizer. <https://code.google.com/p/data-race-test/wiki/ThreadSanitizer>, 2015.
- [7] J. Yang, A. Cui, S. Stolfo, and S. Sethumadhavan. Concurrency attacks. In *the Fourth USENIX Workshop on Hot Topics in Parallelism (HOTPAR '12)*, June 2012.
- [8] W. Zhang, J. Lim, R. Olichandran, J. Scherpelz, G. Jin, S. Lu, and T. Reps. ConSeq: detecting concurrency bugs through sequential errors. In *Sixteenth International Conference on Architecture Support for Programming Languages and Operating Systems (ASPLOS '11)*, pages 251–264, Mar. 2011.
- [9] W. Zhang, C. Sun, and S. Lu. ConMem: detecting severe concurrency bugs through an effect-oriented approach. In *Fifteenth International Conference on Architecture Support for Programming Languages and Operating Systems (ASPLOS '10)*, pages 179–192, Mar. 2010.